

Assignment 0: Arithmetic Expressions

Version 1.0

Due: Wednesday, September 12 at 20:00

The objective of this assignment is to gain some initial hands-on programming experience with Haskell.

1 Simple arithmetic expressions

Consider the following algebraic data type representing simple *arithmetic expressions*:

```
data Exp =
    Cst Integer
  | Add Exp Exp
  | Sub Exp Exp
  | Mul Exp Exp
  | Div Exp Exp
  | Pow Exp Exp
-- ... (additional constructors; see next section)
```

That is, an arithmetic expression is either an (unbounded) integer constant, or one of the following five operations on two subexpressions: addition, subtraction, multiplication, division, or power (aka. exponentiation). Note that there is no provision for including explicit parentheses in the representation, as the tree structure of an `Exp`-typed value already encodes the intended grouping: the arithmetic expression conventionally written as $2 \times (3 + 4)$ is represented as `Mul (Cst 2) (Add (Cst 3) (Cst 4))`, whereas $(2 \times 3) + 4$ corresponds to `Add (Mul (Cst 2) (Cst 3)) (Cst 4)`.

1.1 Printing expressions

Define a function

```
showExp :: Exp -> String
```

that renders an arithmetic expression as a string in ordinary mathematical notation, as in the examples above. Use “+”, “-”, “*”, “/”, and “^” to represent the five arithmetic operators. Insert enough parentheses in the output to ensure that any two different `Exp` trees have distinct renderings. It’s fine to include nominally redundant (according to the usual mathematical conventions) parentheses, e.g., `"((2*3)+4)"`.

If the expression to be printed is not one of the above 6 forms (i.e., belongs to the commented-out part of the listing), your code should explicitly report (using the built-in function `error`) the problem with an appropriate message, rather than crash out with a nonexhaustive pattern-match error.

1.2 Evaluating expressions

Arithmetic expressions can be evaluated to numeric results. In this exercise, we only consider integer arithmetic, with n/m (for $m \neq 0$) defined as $\lfloor \frac{n}{m} \rfloor$, where $\lfloor x \rfloor$ (the *floor* of x) is the greatest integer less than or equal to x . (So $\lfloor \frac{7}{2} \rfloor = 3$, while $\lfloor \frac{-7}{2} \rfloor = -4$.) Also, we require that the exponent (second subexpression) in a `Pow`-operation is non-negative, and we specify that $n^0 = 1$ for all n , including 0.

Define a Haskell function

```
evalSimple :: Exp -> Integer
```

such that `evalSimple e` computes the value of e , under the interpretation of the arithmetic operators specified above. If an error occurs (e.g., division by zero), it is fine to just abort with the relevant Haskell runtime error. And, like for printing, expressions not in the “simple” fragment of `Exp` should be explicitly reported as errors. If there are multiple errors in an expression, it doesn’t matter which one gets reported.

2 Extended arithmetic expressions

We now consider a richer class of expressions, given by the datatype:

```
data Exp =
-- ... (6 constructors from above)
  | If {test, yes, no :: Exp}
  | Var VName
  | Let {var :: VName, aux, body :: Exp}
  | Sum {var :: VName, from, to, body :: Exp}
```

```
type VName = String
```

Here, the expression form `If e_1 e_2 e_3` (or, more verbosely, `If {test = e_1 , yes = e_2 , no = e_3 }`) represents a *conditional expression* (analogous to $e_1 ? e_2 : e_3$ in C). That is, its value is either the value of e_2 , if e_1 evaluates to a non-zero number; or the value of e_3 , if e_1 evaluates to zero. Only the relevant branch is evaluated; for example, evaluating the expression `If (Sub (Cst 2) (Cst 2)) (Div (Cst 3) (Cst 0)) (Cst 5)` should return 5, and *not* abort with a division by zero.

The expression `Var v` , where v is a *variable name* (represented as a Haskell string), returns the current value (as defined below) of the variable v . If the variable has no current value, an error is signaled.

(For simplicity, we do not impose any syntactic restrictions on variable names, so that in principle there may be variables confusingly named “42”, or even “Cst 42”. This does not cause any formal problems, because we maintain a firm distinction between the Haskell expressions `Var "Cst 42" :: Exp` and `Cst 42 :: Exp`.)

Conversely, the expression `Let v e_1 e_2` is used to *bind* v to the value of e_1 for the duration of evaluating e_2 ; afterwards, the previous binding (if any) of v is reinstated. Thus, for example, the expression

```
Let {var = "x", aux = Cst 5,
    body = Add (Let {var = "x", aux = Add (Cst 3) (Cst 4),
                    body = Mul (Var "x") (Var "x")})
          (Var "x")}
```

should evaluate to $(3 + 4)^2 + 5 = 54$.

It is deliberately left *unspecified* (i.e., you as the implementer may decide) whether an error occurring in the auxiliary expression e_1 should be signaled if the bound variable v is not actually used in the body e_2 . For instance,

```
Let "x" (Div (Cst 4) (Cst 0)) (Cst 5)
```

is allowed to evaluate to 5, or abort with a division by zero (but nothing else).

Finally, `Sum v e1 e2 e3` corresponds to the mathematical construction $\sum_{v=e_1}^{e_2} e_3$. That is, it first evaluates e_1 and e_2 to numbers n_1 and n_2 , and then computes the sum of evaluating e_3 , in which v is bound to each of the values $n_1, n_1 + 1, \dots, n_2$ in turn. If $n_1 > n_2$, the sum is defined to be 0.

We keep track of variable bindings in an *environment*, which maps variable names to their values (if any). We represent environments as functions:

```
type Env = VName -> Maybe Integer
```

That is, for an environment $r :: \text{Env}$ and variable $v :: \text{VName}$, the application $r\ v$ returns `Nothing` if v has no binding in r , and `Just n` if v is bound to the integer n . The environment in which all variables are unbound can be written as simply `initEnv = \v -> Nothing`. Normally, an environment will only contain bindings for finitely many variables, but this is not enforced.

Define a function

```
extendEnv :: VName -> Integer -> Env -> Env
```

such that `extendEnv v n r` returns a new environment r' , in which v is bound to n , and all other variables have the same bindings as in r .

Then define a function

```
evalFull :: Exp -> Env -> Integer
```

that evaluates an expression in a given environment. As before, errors should be signaled with `error`, and may now – in addition to undefined arithmetic operations – include attempts to access unbound variables. On the other hand, all expression forms `Exp` should now be covered.

Be sure to specify how you chose to deal with errors in unneeded parts of `Let`-expressions, and why. (Simplicity of implementation is a perfectly acceptable justification.)

3 Returning explicit errors

Promptly aborting evaluation with a Haskell runtime error is a fairly drastic step, and in particular makes it impossible to recover from a conceptually non-fatal problem. A more flexible approach makes the evaluator function return an explicit indication of what went wrong, and lets the caller decide what to do next. Accordingly, we first enumerate some possibilities:

```
ArithError =
  EBadVar Vname    -- unbound variable
| EDivZero         -- attempted division by zero
| ENegPower        -- attempted raising to negative power
| EOther String    -- any other errors, if relevant
```

We can then define a Haskell function,

```
evalErr :: Exp -> Env -> Either ArithError Integer
```

such that `evalErr e r` attempts to evaluate e in environment r , as in `evalFull`, but now returns either an error value (e.g. `Left ENegPower`) or a proper result (e.g. `Right 42`). `evalErr` should *never* cause a Haskell runtime error.

We now specify explicitly that all subexpressions are evaluated left-to-right, so that, e.g., when evaluating `Add e1 e2`, if e_1 returns an error, e_2 should not be evaluated. However, it is still unspecified (meaning: you should choose) whether errors in unused `Let`-bindings should be reported or ignored. Again, justify your choice.

Hint: The code of `evalErr` may become somewhat verbose and repetitive; try to abstract common code snippets into (higher-order) auxiliary functions, so that you only have to write them once. Do not attempt to use Haskell's builtin imprecise-exception facility: it's quite finicky, and will most likely not do what you need.

4 Optional extensions

The problems in the section are a bit more challenging, and hence not mandatory, but still recommended for extra practice. You may do either or both.

Note that your performance on the optional problems will not affect whether you pass the assignment: an incomplete or buggy solution will not drag down an otherwise acceptable solution of the mandatory part, but you will hopefully find the additional feedback useful; conversely, you cannot save an otherwise failing solution of the main problems by successfully solving one of the optional ones.

4.1 Printing with minimal parentheses

Define a variant of the `showExp` function,

```
showCompact :: Exp -> String
```

that prints a simple arithmetic expression using the minimal number of parentheses (and no extra spaces). The function must still be one-to-one, i.e., no two different values of type `Exp` should be rendered into the same string. You should assume that the arithmetic operators have the conventional precedences and associativities (e.g. `Add (Cst 2) (Mul (Cst 3) (Cst 4))` should print as `"2+3*4"`, and `Add (Cst 2) (Add (Cst 3) (Cst 4))` as `"2+(3+4)"`). Note in particular that `"2^3^4"` corresponds to 2^{3^4} , which conventionally means $2^{(3^4)}$, not $(2^3)^4$.

Hint: You will probably want to define `showCompact` in terms of an auxiliary recursive function that takes as parameters both the expression to be printed and some additional data needed to determine whether explicit parentheses are needed.

4.2 Explicitly eager/lazy semantics

In `evalErr`, the exact semantics of `Let`-expressions as left partially unspecified. There are in fact two natural interpretations of a `Let v e1 e2`: the *eager* one, where e_1 is always evaluated first, regardless of whether v is used in e_2 ; and the *lazy* one, where e_1 is only evaluated as and when its value is needed in e_2 .

Accordingly, define two functions

```
evalEager :: Exp -> Env -> Either ArithError Integer
evalLazy  :: Exp -> Env -> Either ArithError Integer
```

implementing the two variants (for all `Lets` occurring in the first argument). For example, the call `evalLazy (Let "x" (Var "y") (Cst 0)) initEnv` should now return `Right 0`, whereas `evalEager (Let "x" (Var "y") (Cst 0)) initEnv` should return `Left (BadVar "y")`.

Hint: If your `evalErr` from already implements precisely one of the two behaviors, you can simply use it directly as the relevant definition. And again, you may want to express one or both of the main evaluation functions in terms of another function with additional and/or differently typed parameters.

5 What to hand in

5.1 Code

Content As always, your code should be appropriately commented. In particular, try to give brief informal specifications for any auxiliary “helper” functions you define, whether locally or globally. On the other hand, avoid trivial comments that just rephrase what the code is already saying. Try to use a consistent indentation style, and avoid lines of over 80 characters.

You may (but shouldn’t need to, for this assignment) import additional functionality from the GHC libraries. Do not use any packages that are not part of the Stack LTS-12.6 distribution. If you use any libraries that are not installed by default, be sure to explain how to compile your code, either in the report or in a README file.

Your code should give no warnings when compiled with `ghc(i) -W`; otherwise, add a comment explaining why any such warning is harmless or irrelevant in each particular instance. If some problem in your code prevents the whole file from compiling at all, be sure to comment out the offending part.

Form Your main code should be placed in the file `Arithmetic.hs` (and, if relevant, in other modules from which `Arithmetic` imports). It should only export the requested functionality. Any tests or examples should be put in a separate module.

We have provided a skeleton `Arithmetic.hs` with stub definitions for all the functions you are expected to define. It is very important that you not change the types of the exported functions, as this will make your code incompatible with our testing framework. Do not remove the bindings for any functions you do not implement; just leave them as `undefined`.

The definitions for this assignment (e.g. type `Exp`) are available in file `Definitions.hs`. You should only import this module, and not directly copy its contents into `Arithmetic`. And again, do not modify anything in `Definitions`.

Your code submission should be packaged as a single file `src.zip` (*not* `.rar`, `.tar.gz`, etc.), which should contain only a single directory `src` (which may itself contain additional subdirectories, if needed). Please take care to include only files that constitute your actual submission: your code and supporting materials (e.g., build/run instructions, tests, etc.), but not obsolete/experimental versions, backups, autosave files, revision-control metadata, and the like.

5.2 Report

In addition to the code, you must submit a short (normally 2–3 pages) report, covering the following two points:

- Document any (relevant) *design* and *implementation* choices you made. This includes, but is not limited to, answering any questions explicitly asked in the assignment text. Focus on high-level aspects and ideas, and motivate *why* you did something, not only *what* you did. It’s rarely appropriate to do a detailed function-by-function code walk-through in the report; non-trivial remarks about how the functions work belong in the code as comments.

- Give a honest, justified *assessment* of the quality of your submitted code, and the degree to which it fulfills the requirements of the assignment (to the best of your understanding and knowledge). Be sure to clearly explain any known or suspected deficiencies.

It is very important that you document on what your assessment is based (e.g., wishful thinking, scattered examples, systematic tests, correctness proofs?). Include any automated tests you did with your source submission, and summarize the results in the report.

Your report submission should be a single PDF file named `report.pdf`, uploaded along with `src.zip`. It should include a listing of your code (but not the already provided auxiliary files) as an appendix.

5.3 General

Detailed upload instructions, in particular regarding the logistics of group submissions, can be found on the Absalon submission page.