

High Speed Hashing for Integers and Strings

Mikkel Thorup

September 19, 2018

Abstract

These notes describe the most efficient hash functions currently known for hashing integers and strings. These modern hash functions are often an order of magnitude faster than those presented in standard text books. They are also simpler to implement, and hence a clear win in practice, but their analysis is harder. Some of the most practical hash functions have only appeared in theory papers, and some of them require combining results from different theory papers. The goal here is to combine the information in lecture-style notes that can be used by theoreticians and practitioners alike, thus making these practical fruits of theory more widely accessible.

1 Hash functions

The concept of truly independent hash functions is extremely useful in the design of randomized algorithms. We have a large universe U of keys, e.g., 64-bit numbers, that we wish to map randomly to a range $[m] = \{0, \dots, m-1\}$ of hash values. A *truly random hash function* $h : U \rightarrow [m]$ assigns an independent uniformly random variable $h(x)$ to each key in x . The function h is thus a $|U|$ -dimensional random variable, picked uniformly at random among all functions from U to $[m]$.

Unfortunately truly random hash functions are idealized objects that cannot be implemented. More precisely, to represent a truly random hash function, we need to store at least $|U| \log_2 m$ bits, and in most applications of hash functions, the whole point in hashing is that the universe is much too large for such a representation (at least not in fast internal memory).

The idea is to let hash functions contain only a small element or seed of randomness so that the hash function is sufficiently random for the desired application, yet so that the seed is small enough that we can store it when first it is fixed. As an example, if p is prime, a random hash function $h : [p] \rightarrow [p] = \{0, \dots, p-1\}$ could be $h(x) = (ax + b) \bmod p$ where a and b are random numbers that together form the random seed describing the function. In these notes we will discuss some basic forms of random hashing that are very efficient to implement, and yet have sufficient randomness for some important applications.

1.1 Definition and properties

Definition 1 A hash function $h : U \rightarrow [m]$ is a random variable in the class of all functions $U \rightarrow [m]$, that is, it consists of a random variable $h(x)$ for each $x \in U$.

For a hash function, we care about roughly three things:

Space The size of the random seed that is necessary to calculate $h(x)$ given x ,

Speed The time it takes to calculate $h(x)$ given x ,

Properties of the random variable.

In the next sections we will mention different desirable properties of the random hash functions, and how to implement them efficiently. First we introduce universal hashing in Section 2, then we introduce strongly universal hashing in Section 3. In both cases, we present very efficient hash function if the keys are 32- or 64-bit integers and the hash values are bit strings. In Section 4 we show how we can efficiently produce hash values in arbitrary integer ranges. In Section 5, we show how to hash keys that are strings. Finally, in Section 6, we briefly mention hash functions that have stronger properties than strong universality.

2 Universal hashing

The concept of universal hashing was introduced by Carter and Wegman in [2]. We wish to generate a random hash function $h : U \rightarrow [m]$ from a key universe U to a set of hash values $[m] = \{0, \dots, m-1\}$. We think of h as a random variable following some distribution over functions $U \rightarrow [m]$. We want h to be *universal* which means that for any given distinct keys $x, y \in U$, when h is picked at random (independently of x and y), we have *low collision probability*:

$$\Pr_h[h(x) = h(y)] \leq 1/m.$$

For many applications, it suffices if for some $c = O(1)$, we have

$$\Pr_h[h(x) = h(y)] \leq c/m.$$

Then h is called *c-universal*.

In this chapter we will first give some concrete applications of universal hashing. Next we will show how to implement universal hashing when the key universe is an integer domain $U = [u] = \{0, \dots, u-1\}$ where the integers fit in a machine word, that is, $u \leq 2^w$ where $w = 64$ is the word length. In later chapters we will show how to make efficient universal hashing for large objects such as vectors and variable length strings.

Exercise 2.1 Is the truly independent hash function $h : U \rightarrow [m]$ universal?

Exercise 2.2 If a hash function $h : U \rightarrow [m]$ has collision probability 0, how large must m be?

Exercise 2.3 Let $u \leq m$. Is the identity function $f(x) = x$ a universal hash function $[u] \rightarrow [m]$?

2.1 Applications

One of the most classic applications of universal hashing is *hash tables with chaining*. We have a set $S \subseteq U$ of keys that we wish to store so that we can find any key from S in expected constant time. Let $n = |S|$ and $m \geq n$. We now pick a universal hash function $h : U \rightarrow [m]$, and then create an array L of m lists/chains so that for $i \in [m]$, $L[i]$ is the list of keys that hash to i . Now to find out if a key $x \in U$ is in S , we only have to check if x is in the list $L[h(x)]$. This takes time proportional to $1 + |L[h(x)]|$ (we add 1 because it takes constant time to look up the list even if turns out to be empty).

Assume that $x \notin S$ and that h is universal. Let $I(y)$ be an indicator variable which is 1 if $h(x) = h(y)$ and 0 otherwise. Then the expected number of elements in $L[h(x)]$ is

$$E_h[|L[h(x)]|] = E_h\left[\sum_{y \in S} I(y)\right] = \sum_{y \in S} E_h[I(y)] = \sum_{y \in S} \Pr_h[h(y) = h(x)] = n/m \leq 1.$$

The second equality uses *linearity of expectation*.

Exercise 2.4 (a) What is the expected number of elements in $L[h(x)]$ if $x \in S$?

(b) What bound do you get if h is only 2-universal?

The idea of hash tables goes back to [8], and hash tables were the prime motivation for the introduction of universal hashing in [2]. For a text book description, see, e.g., [3, §11.2].

A different application is that of assigning a unique *signature* $s(x)$ to each key. Thus we want $s(x) \neq s(y)$ for all distinct keys $x, y \in S$. To get this, we pick a universal hash function $s : U \rightarrow [n^3]$. The probability of an error (collision) is calculated as

$$\Pr_s[\exists \{x, y\} \subseteq S : s(x) = s(y)] \leq \sum_{\{x, y\} \subseteq S} \Pr_s[s(x) = s(y)] = \binom{n}{2} / n^3 < 1/(2n).$$

The first inequality is a *union bound*: that the probability of that at least one of multiple events happen is at most the sum of their probabilities.

The idea of signatures is particularly relevant when the keys are large, e.g., a key could be a whole text document, which then becomes identified by the small signature. This idea could also be used in connection with hash tables, letting the list $L[i]$ store the signatures $s(x)$ of the keys that hash to i , that is, $L[i] = \{s(x) | x \in S, h(x) = i\}$. To check if x is in the table we check if $s(x)$ is in $L[h(x)]$.

Exercise 2.5 With $s : U \rightarrow [n^3]$ and $h : U \rightarrow [n]$ independent universal hash functions, for a given $x \in U \setminus S$, what is the probability of a false positive when we search x , that is, what is the probability that there is a key $y \in S$ such that $h(y) = h(x)$ and $s(y) = s(x)$?

Below we study implementations of universal hashing.

2.2 Multiply-mod-prime

Note that if $m \geq u$, we can just let h be the identity (no randomness needed) so we may assume that $m < u$. We may also assume that $m > 1$; for if $m = 1$, then $[m] = \{0\}$ and then we have the trivial case where all keys hash to 0.

The classic universal hash function from [2] is based on a prime number $p \geq u$. We pick a uniformly random $a \in [p]_+ = \{1, \dots, p-1\}$ and $b \in [p] = \{0, \dots, p-1\}$, and define $h_{a,b} : [u] \rightarrow [m]$ by

$$h_{a,b}(x) = ((ax + b) \bmod p) \bmod m \quad (1)$$

Given any distinct $x, y \in [u] \subseteq [p]$, we want to argue that for random a and b that

$$\Pr_{a \in [p]_+, b \in [p]} [h_{a,b}(x) = h_{a,b}(y)] < 1/m. \quad (2)$$

The strict inequality uses our assumption that $m > 1$. Note that with a truly random hash function into $[m]$, the collision probability is exactly $1/m$, so we are claiming that $h_{a,b}$ has a strictly better collision probability.

In most of our proof, we will consider all $a \in [p]$, including $a = 0$. Ruling out $a = 0$, will only be used in the end to get the tight bound from (2). Ruling out $a = 0$ makes sense because all keys collide when $a = 0$.

We need only one basic fact about primes:

Fact 2.1 *If p is prime and $\alpha, \beta \in [p]_+$ then $\alpha\beta \not\equiv 0 \pmod{p}$.*

Let $x, y \in [p]$, $x \neq y$ be given. For given pair $(a, b) \in [p]^2$, define $(q, r) \in [p]^2$ by

$$(ax + b) \bmod p = q \quad (3)$$

$$(ay + b) \bmod p = r. \quad (4)$$

Lemma 2.2 *Equations (3) and (4) define a 1-1 correspondence between pairs $(a, b) \in [p]^2$ and pairs $(q, r) \in [p]^2$.*

Proof For a given pair $(r, q) \in [p]^2$, we will show that there is at most one pair $(a, b) \in [p]^2$ satisfying (3) and (4). Subtracting (3) from (4) modulo p , we get

$$(ay + b) - (ax + b) \equiv a(y - x) \equiv r - q \pmod{p}, \quad (5)$$

We claim that there is at most one a satisfying (5). Suppose there is another a' satisfying (5). Subtracting the equations with a and a' , we get

$$(a - a')(y - x) \equiv 0 \pmod{p},$$

but since $a - a'$ and $y - x$ are both non-zero modulo p , this contradicts Fact 2.1. There is thus at most one a satisfying (5) for given (q, r) . With this a , we need b to satisfy (3), and this determines b as

$$b = (q - ax) \bmod p. \quad (6)$$

Thus, for each pair $(q, r) \in [p]^2$, there is at most one pair $(a, b) \in [p]^2$ satisfying (3) and (4). On the other hand, (3) and (4) define a unique pair $(q, r) \in [p]^2$ for each pair $(a, b) \in [p]^2$. We have p^2 pairs of each kind, so the correspondence must be 1-1. ■

Since $x \neq y$, by Fact 2.1,

$$r = q \iff a = 0. \quad (7)$$

Thus, when we pick $(a, b) \in [p]_+ \times [p]$, we get $r \neq q$.

Returning to the proof of (2), we get a collision $h_{a,b}(x) = h_{a,b}(y)$ if and only if $q \bmod m = r \bmod m$. Let us fix q and set $i = q \bmod m$. There are at most $\lceil p/m \rceil$ values $r \in [p]$ with $r \bmod m = i$ and one of them is $r = q$. Therefore, the number of $r \in [p] \setminus \{q\}$ with $r \bmod m = i = q \bmod m$ is at most $\lceil p/m \rceil - 1 \leq (p + m - 1)/m - 1 = (p - 1)/m$. However, there are values of $i' \in [m]$ with only $\lfloor p/m \rfloor$ values of $q' \in [p]$ with $q' \bmod m = i'$, and then the number of $r' \in [p] \setminus \{q'\}$ with $r' \bmod m = i' = q' \bmod m$ is $\lfloor p/m \rfloor - 1 < \lceil p/m \rceil - 1$. Summing over all p values of $q \in [p]$, we get that the number of $r \in [p] \setminus \{q\}$ with $r \bmod m = i = q \bmod m$ is strictly less than $p(\lceil p/m \rceil - 1) \leq p(p - 1)/m$. Then our 1-1 correspondence implies that there are strictly less than $p(p - 1)/m$ collision pairs $(a, b) \in [p]_+ \times [p]$. Since each of the $p(p - 1)$ pairs from $[p]_+ \times [p]$ are equally likely, we conclude that the collision probability is strictly below $1/m$, as claimed in (2).

Exercise 2.6 Suppose we for our hash function also consider $a = 0$, that is, for random $(a, b) \in [p]^2$, we define the hash function $h_{a,b} : [p] \rightarrow [m]$ by

$$h_{a,b}(x) = ((ax + b) \bmod p) \bmod m.$$

(a) Show that this function may not be universal.

(b) Prove that it is always 2-universal, that is, for any distinct $x, y \in [p]$,

$$\Pr_{(a,b) \in [p]^2} [h_{a,b}(x) = h_{a,b}(y)] < 2/m.$$

2.2.1 Implementation for 64-bit keys

Let us now consider the implementation of our hashing scheme

$$h(x) = ((ax) \bmod p) \bmod m$$

for the typical case of 64-bit keys in a standard imperative programming language such as C. Let's say the hash values are 20 bits, so we have $u = 2^{64}$ and $m = 2^{20}$.

Since $p > u = 2^{64}$, we generally need to reserve more than 64 bits for $a \in [p]_+$, so the product ax has more than 128 bits. To compute ax , we now have the issue that multiplication of w -bit numbers automatically discards overflow, returning only the w least significant bits of the product. However, we can get the product of 32-bit numbers, representing them as 64-bit numbers, and getting the full 64-bit result. We need at least 6 such 64-bit multiplications to compute ax .

Next issue is, how do we compute $ax \bmod p$? For 64-bit numbers we have a general mod-operation, though it is rather slow, and here we have more than 128 bits.

An idea from [2] is to let p be a Mersenne prime, that is, a prime of the form $2^q - 1$. We can here use the Mersenne prime $2^{89} - 1$. The point in using a Mersenne prime is that

$$x \equiv x \bmod 2^q + \lfloor x/2^q \rfloor \pmod{p}.$$

Exercise 2.7 Assuming a language like C supporting 64-bit multiplication, addition, shifts and bit-wise Boolean operations, but no general mod-operation, sketch the code to compute $((ax + b) \bmod p) \bmod m$ with $p = 2^{89} - 1$ and $m = 2^{20}$.

2.3 Multiply-shift

We shall now turn to a truly practical universal hashing scheme proposed by Dietzfelbinger et al. [6], yet ignored by most text books. It generally addresses hashing from w -bit integers to ℓ -bit integers. We pick a uniformly random odd w -bit integer a , and then we compute $h_a : [2^w] \rightarrow [2^\ell]$, as

$$h_a(x) = \lfloor (ax \bmod 2^w) / 2^{w-\ell} \rfloor \quad (8)$$

This scheme gains an order of magnitude in speed over the scheme from (1), exploiting operations that are fast on standard computers. Numbers are stored as bit strings, with the least significant bit to the right. Integer division by a power of two is thus accomplished by a right shift. For hashing 64-bit integers, we further exploit that 64-bit multiplication automatically discards overflow, which is the same as multiplying modulo 2^{64} . Thus, with $w = 64$, we end up with the following C-code:

```
#include <stdint.h> //defines uint64_t as unsigned 64-bit integer.
uint64_t hash(uint64_t x; uint64_t l; uint64_t a) {
    // hashes x universally into l bits using the random odd seed a.
    return (a*x) >> (64-l);
}
```

This scheme is many times faster and simpler to implement than the standard multiply-mod-prime scheme, but the analysis is more subtle.

It is convenient to think of the bits of a number as indexed with bit 0 the least significant bit. The scheme is simply extracting bits $w - \ell, \dots, w - 1$ from the product ax , as illustrated below.



We will prove that multiply-shift is 2-universal, that is, for $x \neq y$,

$$\Pr_{\text{odd } a \in [2^w]} [h_a(x) = h_a(y)] \leq 2/2^\ell = 2/m. \quad (9)$$

We have $h_a(x) = h_a(y)$ if and only if ax and $ay = ax + a(y - x)$ agree on bits $w - \ell, \dots, w - 1$. This match requires that bits $w - \ell, \dots, w - 1$ of $a(y - x)$ are either all 0s or all 1s. More precisely,


if we get no carry from bits $0, \dots, w - \ell$ when we add $a(y - x)$ to ax , then $h_a(x) = h_a(y)$ exactly when bits $w - \ell, \dots, w - 1$ of $a(y - x)$ are all 0s. On the other hand, if we get a carry 1 from bits $0, \dots, w - \ell$ when we add $a(y - x)$ to ax , then $h_a(x) = h_a(y)$ exactly when bits $w - \ell, \dots, w - 1$ of $a(y - x)$ are all 1s. To prove (9), it thus suffices to prove that the probability that bits $w - \ell, \dots, w - 1$ of $a(y - x)$ are all 0s or all 1s is at most $2/2^\ell$.

We will exploit that any odd number z is relatively prime to any power of two:

Fact 2.3 *If α is odd and $\beta \in [2^q]_+$ then $\alpha\beta \not\equiv 0 \pmod{2^q}$.*

Define b such that $a = 1 + 2b$. Then b is uniformly distributed in $[2^{w-1}]$. Moreover, define z to be the odd number satisfying $(y - x) = z2^i$. Then

$$a(y - x) = z2^i + bz2^{i+1}.$$



Now, we prove that $bz \bmod 2^{w-1}$ must be uniformly distributed in $[2^{w-1}]$. First, note that there is a 1-1 correspondence between the $b \in [2^{w-1}]$ and the products $bz \bmod 2^{w-1}$; for if there were another $b' \in [2^{w-1}]$ with $b'z \equiv bz \pmod{2^{w-1}} \iff z(b' - b) \equiv 0 \pmod{2^{w-1}}$, then this would contradict Fact 2.3 since z is odd. But then the uniform distribution on b implies that $bz \bmod 2^{w-1}$ is uniformly distributed. We conclude that $a(y - x) = z2^i + bz2^{i+1}$ has 0 in bits $0, \dots, i - 1$, 1 in bit i , and a uniform distribution on bits $i + 1, \dots, i + w - 1$.

We have a collision $h_a(x) = h_a(y)$ if ax and $ay = ax + a(y - x)$ are identical on bits $w - \ell, \dots, w - 1$. The two are always different in bit i , so if $i \geq w - \ell$, we have $h_a(x) \neq h_a(y)$ regardless of a . However, if $i < w - \ell$, then because of carries, we could have $h_a(x) = h_a(y)$ if bits $w - \ell, \dots, w - 1$ of $a(y - x)$ are either all 0s, or all 1s. Because of the uniform distribution, either event happens with probability $1/2^\ell$, for a combined probability bounded by $2/2^\ell$. This completes the proof of (9).

Exercise 2.8 *Why is it important that a is odd? Hint: consider the case where x and y differ only in the most significant bit.*

Exercise 2.9 *Does there exist a key x such that $h_a(x)$ is the same regardless of the random odd number a ? If so, can you come up with a real-life application where this is a disadvantage?*

3 Strong universality

We will now consider *strong universality* [11]. For $h : [u] \rightarrow [m]$, we consider *pair-wise events* of the form that for given distinct keys $x, y \in [u]$ and possibly non-distinct hash values $q, r \in [m]$, we have $h(x) = q$ and $h(y) = r$. We say a random hash function $h : [u] \rightarrow [m]$ is *strongly universal* if the probability of every pair-wise event is $1/m^2$. We note that if h is strongly universal, it is also universal since

$$\Pr[h(x) = h(y)] = \sum_{q \in [m]} \Pr[h(x) = q \wedge h(y) = q] = m/m^2 = 1/m.$$

Observation 3.1 *An equivalent definition of strong universality is that each key is hashed uniformly into $[m]$, and that every two distinct keys are hashed independently.*

Proof First assume strong universality and consider distinct keys $x, y \in U$. For any hash value $q \in [m]$, $\Pr[h(x) = q] = \sum_{r \in [m]} \Pr[h(x) = q \wedge h(y) = r] = m/m^2 = 1/m$, so $h(x)$ is uniform in $[m]$, and the same holds for $h(y)$. Moreover, for any hash value $r \in [m]$,

$$\begin{aligned} \Pr[h(x) = q \mid h(y) = r] &= \Pr[h(x) = q \wedge h(y) = r] / \Pr[h(y) = r] \\ &= (1/m^2) / (1/m) = 1/m = \Pr[h(x) = q], \end{aligned}$$

so $h(x)$ is independent of $h(y)$. For the converse direction, when $h(x)$ and $h(y)$ are independent, $\Pr[h(x) = q \wedge h(y) = r] = \Pr[h(x) = q] \cdot \Pr[h(y) = r]$, and when $h(x)$ and $h(y)$ are uniform, $\Pr[h(x) = q] = \Pr[h(y) = r] = 1/m$, so $\Pr[h(x) = q] \cdot \Pr[h(y) = r] = 1/m^2$. ■

Emphasizing the independence, strong universality is also called *2-independence*, as it concerns a pair of *two* events.

Exercise 3.1 *Generalize 2-independence. What is 3-independence? k -independence?*

As for universality, we may accept some relaxed notion of strong universality.

Definition 2 *We say a random hash function $h : U \rightarrow [m]$ is strongly c -universal if*

1. *h is c -uniform, meaning for every $x \in U$ and for every hash value $q \in [m]$, we have $\Pr[h(x) = q] \leq c/m$ and*
2. *every pair of distinct keys hash independently.*

Exercise 3.2 *If h is strongly c -universal, what is an upper bound on the pairwise event probability,*

$$\Pr[h(x) = q \wedge h(y) = r]?$$

Exercise 3.3 *Argue that if $h : U \rightarrow [m]$ is strongly c -universal, then h is also c -universal.*

Exercise 3.4 *Is Multiply-Shift strongly c -universal for any constant c ?*

3.1 Applications

One very important application of strongly universal hashing is *coordinated sampling*, which is crucial to the handling of Big Data and machine learning. The basic idea is that we based on small samples can reason about the similarity of huge sets, e.g., how much they have in common, or how different they are.

First we consider sampling from a single set $A \subseteq U$ using a strongly universal hash function $h : U \rightarrow [m]$ and a threshold $t \in \{0, \dots, m\}$. We now sample x if $h(x) < t$, which by uniformity happens with probability t/m for any x . Let $S_{h,t}(A) = \{x \in A \mid h(x) < t\}$ denote the resulting sample from A . Then, by linearity of expectation, $E[|S_{h,t}(A)|] = |A| \cdot t/m$. Conversely, this means that if we have $S_{h,t}(A)$, then we can estimate $|A|$ as $|S_{h,t}(A)| \cdot m/t$.

We note that the universality from Section 2 does not in general suffice for any kind of sampling. If we, for example, take the multiplication-shift scheme from Section 2.3, then we always have $h(0) = 0$, so 0 will always be sampled if we sample anything, that is, if $t > 0$.

The important application is, however, not the sampling from a single set, but rather the sampling from different sets B and C so that we can later reason about the similarity, estimating the sizes of their union $B \cup C$ and intersection $B \cap C$.

Suppose we for two different sets B and C have found the samples $S_{h,t}(B)$ and $S_{h,t}(C)$. Based on these we can compute the sample of the union as the union of the samples, that is, $S_{h,t}(B \cup C) = S_{h,t}(B) \cup S_{h,t}(C)$. Likewise, we can compute the sample of the intersection as $S_{h,t}(B \cap C) = S_{h,t}(B) \cap S_{h,t}(C)$. We can then estimate the size of the union and intersection multiplying the corresponding sample sizes by m/t .

The crucial point here is that the sampling from different sets can be done in a distributed fashion as long as a fixed h and t is shared, *coordinating the sampling* at all locations. This is used, e.g., in machine learning, where we can store the samples of many different large sets. When a new set comes, we sample it, and compare the sample with the stored samples to estimate which other set it has most in common with. Another cool application of coordinated sampling is on the Internet where all routers can store samples of the packets passing through [7]. If a packet is sampled, it is sampled by all routers that it passes, and this means that we can follow the packets route through the network. If the routers did not use coordinated sampling, the chance that the same packet would be sampled at multiple routers would be very small.

Exercise 3.5 Given $S_{h,t}(B)$ and $S_{h,t}(C)$, how would you estimate the size of the symmetric difference $(B \setminus C) \cup (C \setminus B)$?

Below, in our mathematical reasoning, we only talk about the sample $S_{h,t}(A)$ from a single set A . However, as described above, in many applications, A represent a union $B \cup C$ or intersection $B \cap C$ of different sets B and C .

To get a fixed sampling probability t/m for each $x \in U$, we only need that $h : U \rightarrow [m]$ is uniform. This ensures that the estimate $|S_{h,t}(A)| \cdot m/t$ of $|A|$ is unbiased, that is, $E[|S_{h,t}(A)| \cdot m/t] = |A|$. The reason that we also want the pair-wise independence of strong universality is that we want $|S_{h,t}(A)|$ to be concentrated around its mean $|A| \cdot t/m$ so that we can trust the estimate $|S_{h,t}(A)| \cdot m/t$ of $|A|$. The following lemma implies a good concentration when $|A| \gg m/t$.

Lemma 3.2 Let $X = |S_{h,t}(A)|$ and $\mu = E[|X|] = |A| \cdot t/m$. Then $\text{Var}(X) \leq \mu$ and for any $q > 0$,

$$\Pr[|X - \mu| \geq q\sqrt{\mu}] \leq 1/q^2. \quad (10)$$

Proof For $a \in A$, let X_a be the indicator variable for a being sampled, that is $X_a = [h(a) < t]$. Then $X = \sum_{a \in A} X_a$. Since h is strongly universal, for any distinct keys $a, b \in A$, $h(a)$ and $h(b)$ are independent, but then X_a and X_b are also independent random variables. The X_a are thus pairwise independent, so $\text{Var}[X] = \sum_a \text{Var}[X_a]$. Let $p = t/m$ be the common sampling probability for all keys. We defined $\mu = p|A| = E[X]$ to be the expectation of X . For each X_a , we have $E[X_a] = p$, so $\text{Var}[X_a] = p(1-p)$. It follows that $\text{Var}[X] = \mu(1-p) < \mu$. By definition, the standard deviation of X is $\sigma = \sqrt{\text{Var}[X]}$, and by Chebyshev's inequality (see, e.g., [9, Theorem 3.3]), for any $q > 0$,

$$\Pr[|X - \mu| \geq q\sigma] \leq 1/q^2. \quad (11)$$

This implies (10) since $\sigma \leq \sqrt{\mu}$. ■

Exercise 3.6 Suppose that $|A| = 100,000,000$ and $p = t/m = 1/100$. Then $E[X] = \mu = 1,000,000$. Give an upper bound for the probability that $|X - \mu| \geq 10,000$. These numbers correspond to a 1% sampling rate and a 1% error.

3.2 Multiply-mod-prime

The classic strongly universal hashing scheme is a multiply-mod-prime scheme. For some prime p , uniformly at random we pick $(a, b) \in [p]^2$ and define $h_{a,b} : [p] \rightarrow [p]$ by

$$h_{a,b}(x) = (ax + b) \bmod p. \quad (12)$$

To see that this is strongly universal, consider distinct keys $x, y \in [p]$ and possibly non-distinct hash values $q, r \in [p]$, $h_{a,b}(x) = q$ and $h_{a,b}(y) = r$. This is exactly as in (3) and (4), and by Lemma 2.2, we have a 1-1 correspondence between pairs $(a, b) \in [p] \times [p]$ and pairs $(q, r) \in [p]^2$. Since (a, b) is uniform in $[p]^2$ it follows that (q, r) is uniform in $[p]^2$, hence that the pair-wise event $h_{a,b}(x) = q$ and $h_{a,b}(y) = r$ happens with probability $1/p^2$.

Exercise 3.7 Let $m \leq p$. For random $(a, b) \in [p]^2$, define the hash function $h_{a,b} : [p] \rightarrow [m]$ by

$$h_{a,b}(x) = ((ax + b) \bmod p) \bmod m.$$

(a) Argue $h_{a,b}$ is strongly 2-universal.

(b) In the universal multiply-mod-prime hashing from Section 2, we insisted on $a \neq 0$, but now we consider all $a \in [p]$. Why this difference?

3.3 Multiply-shift

We now present a simple generalization from [4] of the universal multiply-shift scheme from Section 2 that yields strong universality. As a convenient notation, for any bit-string z and integers $j > i \geq 0$, $z[i, j) = z[i, j - 1]$ denotes the number represented by bits $i, \dots, j - 1$ (bit 0 is the least significant bit, which confusingly, happens to be rightmost in the standard representation), so

$$z[i, j) = \lfloor (z \bmod 2^j) / 2^i \rfloor.$$

To hash $[2^w] \rightarrow [2^\ell]$ we may pick any $\bar{w} \geq w + \ell - 1$. For any pair $(a, b) \in [\bar{w}]^2$, we define $h_{a,b} : [2^w] \rightarrow [2^\ell]$ by

$$h_{a,b}(x) = (ax + b)[\bar{w} - \ell, \bar{w}). \quad (13)$$

As for the universal multiply shift, we note that the scheme of (13) is easy to implement with convenient parameter choices, e.g., with $\bar{w} = 64$, $w = 32$ and $\ell = 20$, we get the C-code:

```
#include <stdint.h>
// defines uint32/64_t as unsigned 32/64-bit integer.
uint64_t hash(uint32_t x; uint32_t l; uint64_t a; uint64_t b; ) {
    // hashes x strongly universally into l bits
    // using the random seeds a and b.
    return (a*x+b) >> (64-l);
}
```

We will prove that the scheme from (13) is strongly universal. In the proof we will reason a lot about uniformly distributed variables, e.g., if $X \in [m]$ is uniformly distributed and β is a constant integer, then $(X + \beta) \bmod m$ is also uniformly distributed in $[m]$. More interestingly, we have

Fact 3.3 *Consider two positive integers α and m that are relatively prime, that is, α and m have no common prime factor. If X is uniform in $[m]$, then $(\alpha X) \bmod m$ is also uniformly distributed in $[m]$. Important cases are (a) if $\alpha < m$ and m is prime, and (b) if α is odd and m is a power of two.*

Proof We want to show that for every $y \in [m]$ there is at most one $x \in [m]$ such that $(\alpha x) \bmod m = y$, for then there must be exactly one $x \in [m]$ for each $y \in [m]$, and vice versa. Suppose we had distinct $x_1, x_2 \in [m]$ such that $(\alpha x_1) \bmod m = y = (\alpha x_2) \bmod m$. Then $\alpha(x_2 - x_1) \bmod m = 0$, so m is a divisor of $\alpha(x_2 - x_1)$. By the fundamental theorem of arithmetic, every positive integer has a unique prime factorization, so all prime factors of m have to be factors of $\alpha(x_2 - x_1)$ in same or higher powers. Since m and α are relatively prime, no prime factor of m is factor of α , so the prime factors of m must all be factors of $x_2 - x_1$ in same or higher powers. Therefore m must divide $x_2 - x_1$, contradicting the assumption $x_1 \not\equiv x_2 \pmod{m}$. Thus, as desired, for any $y \in [m]$, there is at most one $x \in [m]$ such that $(\alpha x) \bmod m = y$. ■

Theorem 3.4 *When $a, b \in [2^{\bar{w}}]$ are uniform and independent, the multiply-shift scheme from (13) is strongly universal.*

Proof Consider any distinct keys $x, y \in [2^w]$. We want to show that $h_{a,b}(x)$ and $h_{a,b}(y)$ are independent uniformly distributed variables in $[2^\ell]$.

Let s be the index of the least significant 1-bit in $(y - x)$ and let z be the odd number such that $(y - x) = z2^s$. Since z is odd and a is uniform in $[2^{\overline{w}}]$, by Fact 3.3 (b), we have that az is uniform in $[2^{\overline{w}}]$. Now $a(y - x) = az2^s$ has all 0s in bits $0, \dots, s - 1$ and a uniform distribution on bits $s, \dots, s + \overline{w} - 1$. The latter implies that $a(y - x)[s, \dots, \overline{w} - 1]$ is uniformly distributed in $[2^{\overline{w}-s}]$.

Consider now any fixed value of a . Since b is still uniform in $[2^{\overline{w}}]$, we get that $(ax + b)[0, \overline{w})$ is uniformly distributed, implying that $(ax + b)[s, \overline{w})$ is uniformly distributed. This holds for any fixed value of a , so we conclude that $(ax + b)[s, \overline{w})$ and $a(y - x)[s, \overline{w})$ are independent random variables, each uniformly distributed in $[2^{\overline{w}-s}]$.

Now, since $a(y - x)[0, s) = 0$, we get that

$$(ay + b)[s, \infty) = ((ax + b) + a(y - x))[s, \infty) = (ax + b)[s, \infty) + a(y - x)[s, \infty).$$

The fact that $a(y - x)[s, \overline{w})$ is uniformly distributed independently of $(ax + b)[s, \overline{w})$ now implies that $(ay + b)[s, \overline{w})$ is uniformly distributed independently of $(ax + b)[s, \overline{w})$. However, $\overline{w} \geq w + \ell - 1$ and $s < w$ so $s \leq w - 1 \leq \overline{w} - \ell$. Therefore $h_{a,b}(x) = (ax + b)[\overline{w} - \ell, \overline{w})$ and $h_{a,b}(y) = (ay + b)[\overline{w} - \ell, \overline{w})$ are independent uniformly distributed variables in $[2^\ell]$. ■

In order to reuse the above proof in more complicated settings, we crystallize a technical lemma from the last part:

Lemma 3.5 *Let $\overline{w} \geq w + \ell - 1$. Consider a random function $g : U \rightarrow [2^{\overline{w}}]$ with the property that there for any distinct $x, y \in U$ exists a positive $s < w$, determined by x and y (and not by g), such that $(g(y) - g(x))[0, s) = 0$ while $(g(y) - g(x))[s, \overline{w})$ is uniformly distributed in $[2^{\overline{w}-s}]$. For b uniform in $[2^{\overline{w}}]$ and independent of g , define $h_{g,b} : U \rightarrow [2^\ell]$ by*

$$h_{g,b}(x) = (g(x) + b)[\overline{w} - \ell, \overline{w}).$$

Then $h_{g,b}(x)$ is strongly universal.

In the proof of Theorem 3.4, we would have $U = [2^w]$ and $g(x) = ax[0, \overline{w})$, and s was the least significant set bit in $y - x$.

3.4 Vector multiply-shift

Our strongly universal multiply shift scheme generalizes nicely to vector hashing. The goal is to get strongly universal hashing from $[2^w]^d$ to 2^ℓ . With $\overline{w} \geq w + \ell - 1$, we pick independent uniform $a_0, \dots, a_{d-1}, b \in [2^{\overline{w}}]$ and define $h_{a_0, \dots, a_{d-1}, b} : [2^w]^d \rightarrow [2^\ell]$ by

$$h_{a_0, \dots, a_{d-1}, b}(x_0, \dots, x_{d-1}) = \left(\left(\sum_{i \in [d]} a_i x_i \right) + b \right) [\overline{w} - \ell, \overline{w}). \quad (14)$$

Theorem 3.6 *The vector multiply-shift scheme from (14) is strongly universal.*

Proof We will use Lemma 3.5 to prove that this scheme is strongly universal. We define $g : [2^w]^d \rightarrow [2^{\overline{w}}]$ by

$$g(x_0, \dots, x_{d-1}) = \left(\sum_{i \in [d]} a_i x_i \right) [0, \overline{w}).$$

Consider two distinct keys $x = (x_0, \dots, x_{d-1})$ and $y = (y_0, \dots, y_{d-1})$. Let j be an index such that $x_j \neq y_j$ and such that the index s of the least significant set bit is as small as possible. Thus $y_j - x_j$ has 1 in bit s , and all $i \in [d]$ have $(y_j - x_j)[0, s) = 0$. As required by Lemma 3.5, s is determined from the keys only, as required by Lemma 3.5. Then

$$(g(y) - g(x))[0, s) = \left(\sum_{i \in [d]} a_i (y_i - x_i) \right) [0, s) = 0$$

regardless of a_0, \dots, a_{d-1} . Next we need to show that $(g(y) - g(x))[s, \overline{w})$ is uniformly distributed in $[2^{\overline{w}-s}]$. The trick is to first fix all $a_i, i \neq j$, arbitrarily, and then argue that $(g(y) - g(x))[s, \overline{w})$ is uniform when a_i is uniform in $[2^{\overline{w}}]$. Let z be the odd number such that $z2^s = y_j - x_j$. Also, let Δ be the constant defined by

$$\Delta 2^s = \sum_{i \in [d], i \neq j} a_i (y_i - x_i).$$

Now

$$g(y) - g(x) = (a_j z + \Delta) 2^s.$$

With z odd and Δ a fixed constant, the uniform distribution on $a_j \in [2^{\overline{w}}]$ implies that $(a_j z + \Delta) \bmod 2^{\overline{w}}$ is uniform in $[2^{\overline{w}}]$ but then $(a_j z + \Delta) \bmod 2^{\overline{w}-s} = (g(y) - g(x))[s, \overline{w})$ is also uniform in $[2^{\overline{w}-s}]$. Now Lemma 3.5 implies that the vector multiply-shift scheme from (14) is strongly universal. ■

Exercise 3.8 *Corresponding to the universal hashing from Section 2, suppose we tried with $\overline{w} = w$ and just used random odd $a_0, \dots, a_{d-1} \in [2^w]$ and a random $b \in [2^w]$, and defined*

$$h_{a_0, \dots, a_{d-1}, b}(x_0, \dots, x_{d-1}) = \left(\left(\sum_{i \in [d]} a_i x_i \right) + b \right) [w - \ell, w).$$

Give an instance showing that this simplified vector hashing scheme is not remotely universal.

Our vector hashing can also be used for universality, where it gives collision probability $1/2^\ell$. As a small tuning, we could skip adding b , but then we would only get the same $2/2^\ell$ bound as we had in Section 2.

3.5 Pair-multiply-shift

A cute trick from [1] allows us roughly double the speed of vector hashing, the point being that multiplication is by far the slowest operation involved. We will use exactly the same parameters and seeds as for (14). However, assuming that the dimension d is even, we replace (14) by

$$h_{a_0, \dots, a_{d-1}, b}(x_0, \dots, x_{d-1}) = \left(\left(\sum_{i \in [d/2]} (a_{2i} + x_{2i+1})(a_{2i+1} + x_{2i}) \right) + b \right) [\overline{w} - \ell, \overline{w}]. \quad (15)$$

This scheme handles pairs of coordinates $(2i, 2i + 1)$ with a single multiplication. Thus, with $\overline{w} = 64$ and $w = 32$, we handle each pair of 32-bit keys with a single 64-bit multiplication.

Exercise 3.9 (a bit more challenging) *Prove that the scheme defined by (15) is strongly universal. One option is to prove a tricky generalization of Lemma 3.5 where $(g(y) - g(x))[0, s)$ may not be 0 but can be any deterministic function of x and y . With this generalization, you can make a proof similar to that for Theorem 3.6 with the same definition of j and s .*

Above we have assumed that d is even. In particular this is a case, if we want to hash an array of 64-bit integers, but cast it as an array of 32-bit numbers. If d is odd, we can use the pair-multiplication for the first $\lfloor d/2 \rfloor$ pairs, and then just add $a_d x_d$ to the sum.

4 Fast hashing to arbitrary ranges

Using variants of multiply-shift, we have shown very efficient methods for hashing into ℓ -bit hash values for $\ell \leq 32$, but what if we want hash values in $[m]$ for any given $m < 2^\ell$.

The general problem is if we have a good hash function $h : U \rightarrow [M]$, and now we want hash values in $[m]$ where $m < M$. What we need is a function $r : [M] \rightarrow [m]$ that is *most uniform* in the sense that for any $z \in [m]$, the number of $y \in [M]$ that map to z is either $\lfloor M/m \rfloor$ or $\lceil M/m \rceil$.

Exercise 4.1 *Prove that if h is strongly c -universal and r is most uniform, then $r \circ h$, mapping x to $r(h(x))$ is strongly $(1 + m/M)c$ -universal.*

An example of a most uniform function r is $y \mapsto y \bmod m$. We already used this r in Exercise 3.7 where we first had a strongly 1-universal hash function into $[p]$ and then applied $\bmod m$. However, computing \bmod is much more expensive than a multiplication on most computers. An alternative way to get a most uniform hash function $r : [M] \rightarrow [m]$ is to set

$$r(y) = \lfloor ym/M \rfloor. \quad (16)$$

Exercise 4.2 *Prove that r defined in (16) is most uniform.*

While (16) is not fast in general, it is very fast if $M = 2^\ell$ is a power-of-two, for then the division is just a right shift, and then (16) is computed by $(y*m) >> 1$. One detail to note here is that the product ym has to be computed in full with no discarded overflow, e.g., if y and m are 32-bit integers, we need 64-bit multiplication. Combining this with the code for strongly universal multiply-shift from Section 3.3, we hash a 32-bit integer x to a number in $[m]$, $m < 2^{32}$, using the C-code:

```
#include <stdint.h>
// defines uint32/64_t as unsigned 32/64-bit integer.
uint32_t hash(uint32_t x; uint32_t m; uint64_t a; uint64_t b; ) {
    // hashes x strongly universally into the range [m]
    // using the random seeds a and b.
    return ((a*x+b)>>32)*m>>32;
}
```

Above, x and m are 32-bit integers while a and b are uniformly random 64-bit integers. We note that all the above calculations are done with 64-bit integers since they all involve 64-bit operands. As required in Section 3.3, we automatically discard the overflow beyond 64 bits from $a*x$. However $(a*x+b) >> 32$ only uses the 32 least significant bits, so multiplied with the 32-bit integer m , we get the exact product in 64 bits with no overflow. From the above Exercises, it immediately follows that the C-code function h is a strongly 2-universal hash function from 32-bit integers to integers in $[m]$.

4.1 Hashing to larger ranges

So far, we have been focused on hashing to 32-bit numbers or less. If we want larger hash values, the most efficient method is often just to use multiple hash functions. The idea is captured by the following exercise.

Exercise 4.3 *Let h_0 be a strongly c_0 -universal hash function from U to R_0 and h_1 be a strongly c_1 -universal hash function from U to R_1 . Define the combined hash function $h : U \rightarrow R_0 \times R_1$ by*

$$h(x) = (h_0, h_1).$$

Prove that h is strongly c_0c_1 -universal.

A simple application of Exercise 4.3 is if h_0 and h_1 are the strongly universal hash functions from Section 3.5, generating 32-bit hash values from 64-bit keys. Then the combined hash function h is a strongly universal function from 64-bit keys to 64-bit hash values. It is the fastest such hash function known, and it uses only two 64-bit multiplications.

5 String hashing

5.1 Hashing vector prefixes

Sometimes what we really want is to hash vectors of length up to D but perhaps smaller. As in the multiply-shift hashing schemes, we assume that each coordinate is from $[2^w]$. The simple point is that we only want to spend time proportional to the actual length $d \leq D$. With $\overline{w} \geq w + \ell - 1$, we pick independent uniform $a_0, \dots, a_{D-1} \in [2^{\overline{w}}]$. For even d , we define $h : \bigcup_{\text{even } d \leq D} [2^w]^d \rightarrow [2^\ell]$ by

$$h_{a_0, \dots, a_D}(x_0, \dots, x_{d-1}) = \left(\left(\sum_{i \in [d/2]} (a_{2i} + x_{2i+1})(a_{2i+1} + x_{2i}) \right) + a_d \right) [2^{\overline{w}} - \ell, \overline{w}). \quad (17)$$

Exercise 5.1 *Prove that the above even prefix version of pair-multiply-shift is strongly universal. In the proof you may assume that the original pair-multiply-shift from (15) is strongly universal, as you may have proved in Exercise 3.9. Thus we are considering two vectors $x = (x_0, \dots, x_{d-1})$ and $y = (y_0, \dots, y_{d'-1})$. You should consider both the case $d' = d$ and $d' \neq d$.*

5.2 Hashing bounded length strings

Suppose now that we want to hash strings of 8-bit characters, e.g., these could be the words in a book. Then the nil-character is not used in any of the strings. Suppose that we only want to handle strings up to some maximal length, say, 256.

With the prefix-pair-multiply-shift scheme from (17), we have a very fast way of hashing strings of d 64-bit integers, casting them as $2d$ 32-bit integers. A simple trick now is to allocate a single array x of $256/8 = 32$ 64-bit integers. When we want to hash a string s with c characters, we first set $d = \lceil c/8 \rceil$ (done fast by $d = (c+7) \gg 3$). Next we set $x_{d-1} = 0$, and finally we do a memory copy of s into x (using a statement like `memcpy(x, s, c)`). Finally, we apply (17) to x .

Note that we use the same variable array x every time we want to hash a string s . Let s^* be the image of s created as a $c^* = \lceil c/8 \rceil$ length prefix of x .

Exercise 5.2 *Prove that if s and t are two strings of length at most 256, neither containing the nil-character, then their images s^* and t^* are different. Conclude that we now have a strongly universal hash functions for such strings.*

Exercise 5.3 *Implement the above hash function for strings. Use it in a chaining based hash table, and apply it to count the number of distinct words in a text (take any pdf-file and convert it to ASCII, e.g., using `pdf2txt`).*

To get the random numbers defining your hash functions, you can go to random.org.

One issue to consider when you implement a hash table is that you want the number m of entries in the hash array to be as big as the number of elements (distinct words), which in our case

is not known in advance. Using a hash table of some start size m , you can maintain a count of the distinct words seen so far, and then double the size when the count reaches, say, $m/2$.

Many ideas can be explored for optimization, e.g., if we are willing to accept a small false-positive probability, we can replace each word with a 32- or 64-bit hash value, saying that a word is new only if it has a new hash value.

Experiment with some different texts: different languages, and different lengths. What happens with the vocabulary?

The idea now is to check how much time is spent on the actual hashing, as compared with the real code that both does the hashing and follows the chains in the hash array. However, if we just compute the hash values, and don't use them, then some optimizing compilers, will notice, and just do nothing. You should therefore add up all the hash values, and output the result, just to force the compiler to do the computation.

5.3 Hashing variable length strings

We now consider the hashing of a variable length string $x_0x_1 \cdots x_d$ where all characters belong to some domain $[u]$.

We use the method from [5], which first picks a prime $p \geq u$. The idea is to view x_0, \dots, x_d as coefficients of a degree d polynomial

$$P_{x_0, \dots, x_d}(\alpha) = \sum_{i=0}^d x_i \alpha^i \bmod p$$

over \mathbb{Z}_p . As seed for our hash function, we pick an argument $a \in [p]$, and compute the hash function

$$h_a(x_0 \cdots x_d) = P_{x_0, \dots, x_d}(a).$$

Consider some other string $y = y_0y_1 \cdots y_{d'}, d' \leq d$. We claim that

$$\Pr_{a \in [p]} [h_a(x_0 \cdots x_d) = h_a(y_0 \cdots y_{d'})] \leq d/p$$

The proof is very simple. By definition, the collision happens only if a is root in the polynomial $P_{y_0, \dots, y_{d'}} - P_{x_0, \dots, x_d}$. Since the strings are different, this polynomial is not the constant zero. Moreover, its degree is at most d . Since the degree is at most d , the fundamental theorem of algebra tells us that it has at most d distinct roots, and the probability that a random $a \in [p]$ is among these roots is at most d/p .

Now, for a fast implementation using Horner's rule, it is better to reverse the order of the coefficients, and instead use the polynomial

$$P_{x_0, \dots, x_d}(a) = \sum_{i=0}^d x_{d-i} a^i \bmod p$$

Then we compute $P_{x_0, \dots, x_d}(a)$ using the recurrence

- $H_a^0 = x_0$
- $H_a^i = (aH_a^{i-1} + x_i) \bmod p$
- $P_{x_0, \dots, x_d}(a) = H_a^d$.

With this recurrence, we can easily update the hash value if new character x_{d+1} is added to the end of the string x_{d+1} . It only takes an addition and a multiplication modulo p . For speed, we would let p be a Mersenne prime, e.g. $2^{89} - 1$.

The collision probability d/p may seem fairly large, but assume that we only want hash values in the range $m \leq p/d$, e.g. for $m = 2^{32}$ and $p = 2^{89} - 1$, this would allow for strings of length up to 2^{57} , which is big enough for most practical purposes. Then it suffices to compose the string hashing with a universal hash function from $[p]$ to $[m]$. Composing with the previous multiply-mod-prime scheme, we end up using three random seeds $a, b, c \in [p]$, and then compute the hash function as

$$h_{a,b,c}(x_0, \dots, x_d) = \left(\left(a \left(\sum_{i=0}^d x_{d-i} c^i \right) + b \right) \bmod p \right) \bmod m.$$

Exercise 5.4 Consider two strings \vec{x} and \vec{y} , each of length at most p/m . Argue that the collision probability $\Pr[h_{a,b,c}(\vec{x}) = h(\vec{y})] \leq 2/m$. Thus conclude that $h_{a,b,c}$ is a 2-universal hash function mapping strings of length at most p/m to $[m]$.

Above we can let u be any value bounded by p . With $p = 2^{89} - 1$, we could use $u = 2^{64}$ thus dividing the string into 64-bit characters.

Exercise 5.5 Implement the above scheme and run it to get a 32-bit signature of a book.

Major speed-up The above code is slow because of the multiplications modulo Mersenne primes, one for every 64 bits in the string.

An idea for a major speed up is to divide your string into chunks X_0, \dots, X_j of 32 integers of 64 bits, the last chunk possibly being shorter. We want a single universal hash function $r : \bigcup_{d \leq 32} [2^{64}]^d \rightarrow [2^{64}]$. A good choice would be to use our strongly universal pair-multiply-shift scheme from (17). It only outputs 32-bit numbers, but if we use two different such functions, we can concatenate their hash values in a single 64-bit number.

Exercise 5.6 Prove that if r has collision probability P , and if $(X_0, \dots, X_j) \neq (Y_0, \dots, Y_j)$, then

$$\Pr[(r(X_0), \dots, r(X_j)) = (r(Y_0), \dots, r(Y_j))] \leq P.$$

The point above is that in the above is that $r(X_0), \dots, r(X_j)$ is 32 times shorter than X_0, \dots, X_j . We can now apply our slow variable length hashing based on Mersenne primes to the reduced string $r(X_0), \dots, r(X_j)$. This only adds P to the overall collision probability.

Exercise 5.7 Implement the above tuning. How much faster is your hashing now?

Splitting between short and long strings When implementing a generic string hashing code, we do not know in advance if it is going to be applied mostly to short or to long strings. Our scheme for bounded length strings from Section 5.2 is faster than the generic scheme presented above for variable length strings. In practice it is a good idea to implement both: have the scheme from Section 5.2 implemented for strings of length up to some d , e.g., d could be 32 64-bit integers as in the above blocks, and then only apply the generic scheme if the length is above d .

Major open problem Can we get something simple and fast like multiply-shift to work directly for strings, so that we do not need to compute polynomials over prime fields?

6 Beyond strong universality

In this note, we have focused on universal and strongly universal hashing. However, there are more advanced algorithmic applications that need more powerful hash functions. This led Carter and Wegman [11] to introduce k -independent hash functions. A hash random function $H : U \rightarrow [m]$ is k -independent if for any distinct keys $x_1, \dots, x_k \in [u]$, the hash values $H(x_1), \dots, H(x_k)$ are independent random variables, each uniformly distributed in $[m]$. In this terminology, 2-independence is the same as strongly universal. For prime p , we can implement a k -independent $H : [p] \rightarrow [p]$ using k random coefficients $a_0, \dots, a_{k-1} \in [p]$, defining

$$H(x) = \sum_{i=0}^{k-1} a_i x^i \bmod p.$$

However, there is no efficient implementation of k -independent hashing on complex objects such as variable length strings. What we can do instead is to use the signature idea from Section 2, and first hash the complex keys to unique hash values from a limited integer domain $[u]$.

Exercise 6.1 Let $h : U \rightarrow [m]$ map $S \subseteq U$ collision free to $[u]$ and let $H : [u] \rightarrow [m]$ be k -independent. Argue that $H \circ h : U \rightarrow [m]$, mapping x to $H(h(x))$, is a k -independent when restricted to keys in S .

Next we have to design the hash function h so that it is collision free with high enough probability.

Exercise 6.2 Suppose we know that we are only going to deal with a (yet unknown) set S of at most n strings, each of length at most n , set the parameters p and m of $h_{a,b,c}$ in Exercise 5.4 so that the probability that we get any collision between strings in S is at most $1/n^2$.

The idea of using hashing to map a complex domain down to a polynomially sized integer universe, hoping that it is collision free, is referred to as *universe reduction*. This explains why we for more complex hash function can often assume that the universe is polynomially bounded.

We note that k -independent hash functions become quite slow when k is large. Often a more efficient alternative is the tabulation based method surveyed in [10].

Acknowledgment

I would like to thank Jacob Holm and Eva Rotenberg for coming with many good comments and proposals for the text, including several exercises.

References

- [1] J. Black, S. Halevi, H. Krawczyk, T. Krovetz, and P. Rogaway. UMAC: fast and secure message authentication. In *Proc. 19th CRYPTO*, pages 216–233, 1999.
- [2] J. Carter and M. Wegman. Universal classes of hash functions. *J. Comp. Syst. Sci.*, 18:143–154, 1979. Announced at STOC’77.
- [3] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. *Introduction to algorithms*. MIT Press, McGraw-Hill, 3 edition, 2009.
- [4] M. Dietzfelbinger. Universal hashing and k -wise independent random variables via integer arithmetic without primes. In *Proc. 13th STACS, LNCS 1046*, pages 569–580, 1996.
- [5] M. Dietzfelbinger, J. Gil, Y. Matias, and N. Pippenger. Polynomial hash functions are reliable (extended abstract). In *Proc. 19th ICALP, LNCS 623*, pages 235–246, 1992.
- [6] M. Dietzfelbinger, T. Hagerup, J. Katajainen, and M. Penttonen. A reliable randomized algorithm for the closest-pair problem. *J. Algorithms*, 25:19–51, 1997.
- [7] N. G. Duffield and M. Grossglauser. Trajectory sampling with unreliable reporting. *IEEE/ACM Trans. Netw.*, 16(1):37–50, 2008.
- [8] A. I. Dumey. Indexing for rapid random access memory systems. *Computers and Automation*, 5(12):6–9, 1956.
- [9] R. Motwani and P. Raghavan. *Randomized Algorithms*. Cambridge University Press, 1995.
- [10] M. Thorup. Fast and powerful hashing using tabulation (research highlight). *Comm. ACM*, 60(7):94–101, 2017. See also <http://arxiv.org/abs/1505.01523>.
- [11] M. Wegman and J. Carter. New hash functions and their use in authentication and set equality. *J. Comp. Syst. Sci.*, 22:265–279, 1981.